

# ECE 414/517: Project #2

**Due date:** Oct 31, 2024 at 23:59PM EST  
**Instructor:** Fei Liu ([fliu33@utk.edu](mailto:fliu33@utk.edu))  
**TA:** Anik Muhury ([amuhury@vols.utk.edu](mailto:amuhury@vols.utk.edu))

## 1 Deep Q-Networks (DQN) (15 pts writeup)

All questions in the section pertain to DQN. The pseudocode for DQN is provided below.

---

**Algorithm 1** Deep Q-Network (DQN)

---

```
1: Initialize replay buffer  $\mathcal{D}$ 
2: Initialize action-value function  $Q$  with random weights  $\theta$ 
3: Initialize target action-value function  $\hat{Q}$  with weights  $\theta^- = \theta$ 
4: for episode = 1,  $M$  do
5:   Receive initial state  $s_1$ 
6:   for  $t = 1, T$  do
7:     With probability  $\epsilon$  select a random action  $a_t$ 
8:     otherwise select  $a_t = \max_a Q(s_t, a; \theta)$ 
9:     Execute action  $a_t$  and observe reward  $r_t$  and state  $s_{t+1}$ 
10:    Store transition  $(s_t, a_t, r_t, s_{t+1})$  in  $\mathcal{D}$ 
11:    Sample random minibatch  $B$  of transitions from  $\mathcal{D}$ 
12:    for each transition  $(s_j, a_j, r_j, s_{j+1})$  in  $B$  do
13:      if  $s_{j+1}$  is terminal then
14:        Set  $y_j = r_j$ 
15:      else
16:        Set  $y_j = r_j + \gamma \max_{a'} \hat{Q}(s_{j+1}, a'; \theta^-)$ 
17:      end if
18:      Perform gradient descent step on  $(y_j - Q(s_j, a_j; \theta))^2$  with respect to network parameters  $\theta$ 
19:    end for
20:    Every  $C$  steps reset  $\hat{Q} = Q$  by setting  $\theta^- = \theta$ 
21:  end for
22: end for
```

---

In this pseudocode:

- $\mathcal{D}$  is the replay memory which stores transitions.
- $\theta$  are the weights of the Q-network, which are adjusted during training.
- $\theta^-$  are the weights of the target network, which are periodically updated to match  $\theta$ .
- $M$  is the number of episodes over which the training occurs.
- $T$  is the maximum number of steps in each episode.
- $\epsilon$  is the exploration rate, which is typically decayed over time.
- $\gamma$  is the discount factor, used to weigh future rewards.
- $C$  is the frequency with which to update the target network's weights.

### 1.1 Written Questions (15 pts)

- (a) (5 pts) (**written**) What are three key differences between the DQN and  $Q$ -learning algorithms?
- (b) (5 pts) (**written**) When using DQN with a deep neural network, which of the above components would you hypothesize contributes most to performance gains? Justify your answer.
- (c) (5 pts) (**written**) In DQN, the choice of target network update frequency is important. What might happen if the target network is updated every  $10^{15}$  steps for an agent learning to play a simple Atari game like Pong?

## 2 Policy Gradient Methods (50 pts coding + 35 pts writeup)

The goal of this problem is to experiment with policy gradient and its variants, including variance reduction and off-policy methods. Your goals will be to set up policy gradient for both continuous and discrete environments, use a neural network baseline for variance reduction, and implement the off-policy Proximal Policy Optimization algorithm. The starter code has detailed instructions for each coding task and includes a README with instructions to set up your environment. Below, we provide an overview of the key steps of the algorithm.

### 2.1 REINFORCE

Recall the policy gradient theorem,

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(a|s) Q^{\pi_{\theta}}(s, a)]$$

REINFORCE is a Monte Carlo policy gradient algorithm, so we will be using the sampled returns  $G_t$  as unbiased estimates of  $Q^{\pi_{\theta}}(s, a)$ . The REINFORCE estimator can be expressed as the gradient of the following objective function:

$$J(\theta) = \frac{1}{\sum T_i} \sum_{i=1}^{|D|} \sum_{t=1}^{T_i} \log(\pi_{\theta}(a_t^i | s_t^i)) G_t^i$$

where  $D$  is the set of all trajectories collected by policy  $\pi_{\theta}$ , and  $\tau^i = (s_0^i, a_0^i, r_0^i, s_1^i, \dots, s_{T_i}^i, a_{T_i}^i, r_{T_i}^i)$  is trajectory  $i$ .

### 2.2 Baseline

One difficulty of training with the REINFORCE algorithm is that the Monte Carlo sampled return(s)  $G_t$  can have high variance. To reduce variance, we subtract a baseline  $b_{\phi}(s)$  from the estimated returns when computing the policy gradient. A good baseline is the state value function,  $V^{\pi_{\theta}}(s)$ , which requires a training update to  $\phi$  to minimize the following mean-squared error loss:

$$L_{\text{MSE}}(\phi) = \frac{1}{\sum T_i} \sum_{i=1}^{|D|} \sum_{t=1}^{T_i} (b_{\phi}(s_t^i) - G_t^i)^2$$

### 2.3 Advantage Normalization

After subtracting the baseline, we get the following new objective function:

$$J(\theta) = \frac{1}{\sum T_i} \sum_{i=1}^{|D|} \sum_{t=1}^{T_i} \log(\pi_{\theta}(a_t^i | s_t^i)) \hat{A}_t^i$$

where

$$\hat{A}_t^i = G_t^i - b_{\phi}(s_t^i)$$

A second variance reduction technique is to normalize the computed advantages,  $\hat{A}_t^i$ , so that they have mean 0 and standard deviation 1. From a theoretical perspective, we can consider centering the advantages to be simply adjusting the advantages by a constant baseline, which does not change the policy gradient. Likewise, rescaling the advantages effectively changes the learning rate by a factor of  $1/\sigma$ , where  $\sigma$  is the standard deviation of the empirical advantages.

## 2.4 Proximal Policy Optimization

One might notice that the REINFORCE algorithm above (with or without a baseline function) is an on-policy algorithm; that is, we collect some number of trajectories under the current policy network parameters, use that data to perform a single batched policy gradient update, and then proceed to discard that data and repeat the same steps using the newly updated policy parameters. This is in stark contrast to an algorithm like DQN which stores all experiences collected over several past episodes. One might imagine that it could be useful to have a policy gradient algorithm “squeeze” a little more information out of each batch of trajectories sampled from the environment. Unfortunately, while the  $Q$ -learning update immediately allows for this, our derived REINFORCE estimator does not in its standard form.

Ideally, an off-policy policy gradient algorithm will allow us to do multiple parameter updates on the same batch of trajectory data. To get a suitable objective function that allows for this, we need to correct for the mismatch between the policy under which the data was collected and the policy being optimized with that data. Proximal Policy Optimization (PPO) restricts the magnitude of each update to the policy (i.e., through gradient descent) by ensuring the ratio of the current and former policies on the current batch is not too different. In doing so, PPO tries to prevent updates that are “too large” due to the off-policy data, which may lead to performance degradation. This technique is related to the idea of importance sampling which we will examine in detail later in the course. Consider the following ratio  $z_\theta$ , which measures the probability ratio between a current policy  $\pi_\theta$  (the “actor”) and an old policy  $\pi_\theta^{\text{old}}$ :

$$z_\theta(s_t^i, a_t^i) = \frac{\pi_\theta(a_t^i | s_t^i)}{\pi_{\theta^{\text{old}}}(a_t^i | s_t^i)}$$

To do so, we introduce the clipped PPO loss function, shown below, where  $\text{clip}(x, a, b)$  outputs  $x$  if  $a \leq x \leq b$ ,  $a$  if  $x < a$ , and  $b$  if  $x > b$ :

$$J_{\text{clip}}(\theta) = \frac{1}{\sum T_i} \sum_{i=1}^{|D|} \sum_{t=1}^{T_i} \min(z_\theta(s_t^i, a_t^i) \hat{A}_t^i, \text{clip}(z_\theta(s_t^i, a_t^i), 1 - \epsilon, 1 + \epsilon) \hat{A}_t^i)$$

where  $\hat{A}_t^i = G_t^i - V_\phi(s_t^i)$ . Note that in this context, we will refer to  $V_\phi(s_t^i)$  as a “critic”; we will train this like the baseline network described above.

To train the policy, we collect data in the environment using  $\pi_\theta^{\text{old}}$  and apply gradient ascent on  $J_{\text{clip}}(\theta)$  for each update. After every  $K$  updates to parameters  $[\pi, \phi]$ , we update the old policy  $\pi_\theta^{\text{old}}$  to equal  $\pi_\theta$ .

## 2.5 Coding Questions (50 pts)

The functions that you need to implement in `network_utils.py`, `policy.py`, `policy_gradient.py`, and `baseline_network.py` are enumerated here. Detailed instructions for each function can be found in the comments in each of these files.

Note: The “batch size” for all the arguments is  $\sum T_i$  since we already flattened out all the episode observations, actions, and rewards for you.

In `network_utils.py`, you need to implement:

- `build_mlp`

In `policy.py`, you need to implement:

- `BasePolicy.act`
- `CategoricalPolicy.action_distribution`
- `GaussianPolicy.__init__`
- `GaussianPolicy.std`

- `GaussianPolicy.action_distribution`

In `policy_gradient.py`, you need to implement:

- `PolicyGradient.init_policy`
- `PolicyGradient.get_returns`
- `PolicyGradient.normalize_advantage`
- `PolicyGradient.update_policy`

In `baseline_network.py`, you need to implement:

- `BaselineNetwork.__init__`
- `BaselineNetwork.forward`
- `BaselineNetwork.calculate_advantage`
- `BaselineNetwork.update_baseline`

In `ppo.py`, you need to implement:

- `PPO.update_policy`

## 2.6 Debugging

To help debug and verify that your implementation is correct, we provide a set of sanity checks below that pass with a correct implementation. Note that these are not exhaustive (i.e., they do not verify that your implementation is correct) and that you may notice oscillation of the average reward across training.

Across most seeds:

- Policy gradient (without baseline) on Pendulum should achieve around an average reward of 100 by iteration 10.
- Policy gradient (with baseline) on Pendulum should achieve around an average reward of 700 by iteration 20.
- PPO on Pendulum should achieve an average reward of 200 by iteration 20.
- All methods should reach an average reward of 200 on Cartpole, 1000 on Pendulum, and 200 on Cheetah at some point.

## 2.7 Writeup Questions (35 pts)

- (a) (5 pts) To compute the REINFORCE estimator, you will need to calculate the values  $\{G_t\}_{t=1}^T$  (we drop the trajectory index  $i$  for simplicity), where

$$G_t = \sum_{t'=t}^T \gamma^{t'-t} r_{t'}$$

Naively, computing all these values takes  $O(T^2)$  time. Describe how to compute them in  $O(T)$  time.

- (b) (5 pts) Consider the cases in the gradient of the clipped PPO loss function equals 0.

Express these cases mathematically and explain why PPO behaves in this manner.

- (c) (5 pts) Notice that the method which samples actions from the policy also returns the log-probability with which the sampled action was taken. Why does REINFORCE not need to cache this information while PPO does? Suppose this log-probability information had not been collected during the rollout. How would that affect the implementation (that is, change the code you would write) of the PPO update?
- (d) (20 pts) The general form for running your policy gradient implementation is as follows:

```
python main.py --env-name ENV --seed SEED --METHOD
```

ENV should be `cartpole`, `pendulum`, or `cheetah`, METHOD should be either `baseline`, `no-baseline`, or `ppo`, and SEED should be a positive integer.

For the `cartpole` and `pendulum` environments, we will consider 3 seeds (`seed = 1, 2, 3`). For `cheetah`, we will only require one seed (`seed = 1`) since it's more computationally expensive, but we strongly encourage you to run multiple seeds if you are able to. Run each of the algorithms we implemented (PPO, PG with baseline, PG without baseline) across each seed and environment. In total, you should end up with at least 21 runs.

Plot the results using:

```
python plot.py --env-name ENV --seeds SEEDS
```

where SEEDS should be a comma-separated list of seeds which you want to plot (e.g. `-seeds 1,2,3`). **Please include the plots (one for each environment) in your writeup, and comment on the performance of each method.**

We have the following expectations about performance to receive full credit:

- `cartpole`: Should reach the max reward of 200 (although it may not stay there)
- `pendulum`: Should reach the max reward of 1000 (although it may not stay there)
- `cheetah`: Should reach at least 200 (could be as large as 900)